


Running a 1.84M Generative Micro-LM on the ESP32-S3

Bare-Metal C++ Forward Pass • ~12-13 tokens/sec



TARGET SILICON

ESP32-S3-WROOM-1

Dual-Core Xtensa LX7 @ 240MHz

OCTAL MEMORY

16MB Flash | 8MB PSRAM

80MHz High-Speed Memory Bus

QUANTIZED PAYLOAD

1.79 MB INT8 Model

Directly mapped in Octal PSRAM

SYSTEM FOOTPRINT

Zero Runtime Overhead

Pure C++ without TFLite or ONNX

Companion Robot Intelligence

THE GOAL

On-Device Generative AI

Executing an autoregressive transformer locally on microcontroller silicon.

THE PERSONA

Meet Kibo (希望 - Hope)

Lightweight conversational companion robot persona running on-chip.

THE FRAMEWORK

ESP32 Micro-LM Engine

A lightweight, modular, and extensible C++ inference engine.

ESP32-S3 Hardware Limits

COMPUTE ENGINE

240 MHz — Xtensa Dual-Core LX7

- SIMD-aligned vector dot-products
- Deterministic cycle execution without garbage collection

MEMORY BUDGET

8 MB — Octal PSRAM (80MHz)

- 1.79 MB Model Weights (INT8)
- Dynamic 128-token Key-Value (KV) working cache

INTERNAL SRAM

380 KB — Internal SRAM Pool

- ~75 KB runtime heap allocation
- Leaves 300+ KB free for peripherals and sensors

End-to-End Pipeline Diagram

Workflow from Host Training to On-Chip Microcontroller Execution:



Model Quantization Benchmark

Evaluation across 4 precision formats on 1.84M parameters:

Format	Model Size	Compression	Cosine Sim	Status
FP32	7.02 MB	0.0% (Base)	100.00%	High PSRAM load
FP16	3.51 MB	50.0%	100.00%	Software emulated
INT8 (Selected)	1.79 MB	74.4%	100.00%	Selected (Optimal)
INT4	0.88 MB	87.2%	98.87%	+26% unpack latency

KEY BENCHMARK TAKEAWAY

INT8 Delivers the Optimal Balance for ESP32-S3

- INT8 reduces PSRAM footprint by 74.4% while preserving 100.00% FP32 output quality.
- Byte-aligned SIMD execution provides maximum inference throughput (~12-13 tok/s).

```
serial_chat.py - 115200 baud

/ import threading
8

Problems Output Debug Console Terminal Ports ESP-IDF
(bash) aiot@aiot-Legion-T7-DLCB-03:~/Projects/KIBO-MICROLMS ./chat kibo_esp32_live.sh
kibo-client: connected to /dev/ttyUSB0 @ 115200 baud (Ctrl+C to exit)

ESP-ROM:esp32s3-20210327
Build:Mar 27 2021
rst:0x1 (POWERON),boot:0x28 (SPI_FAST_FLASH_BOOT)
SPIWP:0xee
mode:DIO, clock div:1
load:0x3fce2820,len:0xfbf0
load:0x403c8700,len:0xaf4
load:0x403cb700,len:0x2e24
entry 0x403c8898

kibo-mcu [1.84M params | int8 | xtensa-lx7 @ 240MHz]
heap: 313524 bytes | psram: 8384788 bytes
[kibo] working buffers allocated (302116 bytes free heap)
[kibo] psram buffer loaded at 0x3c2d4ef4 (1.79 MB)
[kibo] tensor layout: vocab=91 embd=192 layers=4 tensors=53
[kibo] int8 inference engine initialized
ready. enter prompt (e.g. 'hello kibo', 'what is 25 times 4')

User: hello

User: hello
Kibo: [HAPPY] Hello there! Great to see you! How can Kibo help you today?

> [63 tokens generated in 5.17 s | Speed: 12.2 tok/s]

User: how are you?

User: how are you?
Kibo: [HAPPY] I am feeling energetic and super ready to help you!

> [55 tokens generated in 4.53 s | Speed: 12.2 tok/s]

User: what is 100 times 7?

User: what is 100 times 7?
[Tool CALC: 100 * 7 = 700]
```

Physical chip verified

8MB Octal PSRAM

Zero quantization loss

Validated on 3 Form Factors



ESP32-S3 DevKit N16R8

16MB Flash | 8MB Octal PSRAM

✓ 100% Tested & Verified on Hardware



Arduino Nano ESP32

16MB Flash | 8MB Octal PSRAM

✓ 100% Tested & Verified on Hardware



Seeed Studio XIAO S3

8MB Flash | 8MB Octal PSRAM

✓ 100% Tested & Verified on Hardware

100% Open Source on GitHub (MIT License)

github.com/fitranurmayadi/esp32-microlm

Clone, train your custom dataset, and flash to your ESP32!